

You Only Look Once: Unified, Real-Time Object Detection

Joseph Redmon, Santosh Divvala, Ross Girshick, Ali Farhadi

Presented by:

Edward Kwao

Department of Electrical and Computer Engineering

April 23, 2025

Contents

❑ Part 1: YOLO Background

- I. Introduction
- II. YOLO Unified Detection
- III. YOLO Network Design

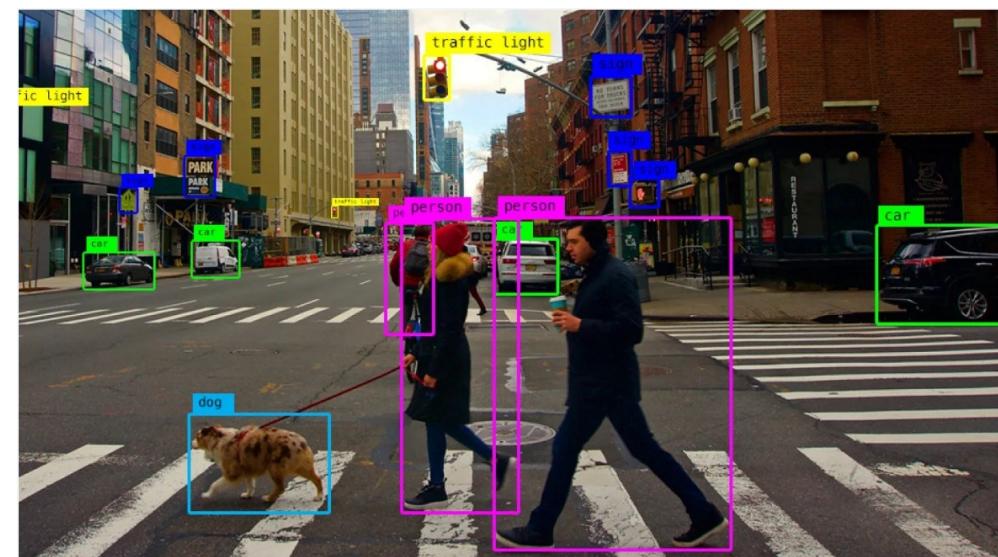
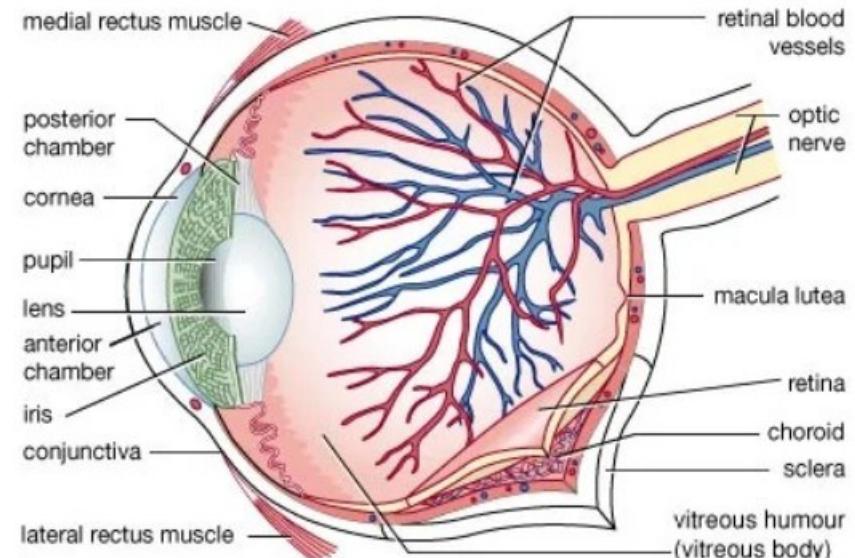
❑ Part 2: YOLO Simulation

- I. Training & Testing YOLO Using Darknet Framework
- II. YOLO Results
- III. YOLO Limitations
- IV. YOLO vs Other Object Detection Systems

❑ Part 3: Conclusion

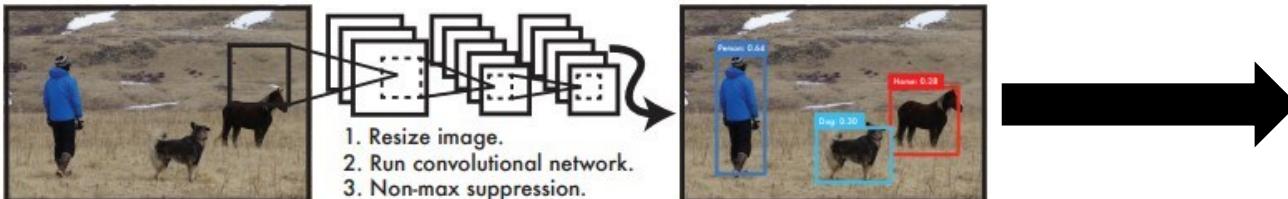
Introduction (1/3)

- ❑ **The human visual system is both fast and accurate**
 - Enables us to handle complex activities like driving with little conscious thought or navigating dark environments.
 - ❑ To replicate this capability, we need object detection algorithms that are also fast and accurate
 - To support autonomous driving.
 - To support assistive technologies that provide real-time visual feedback to users.
 - ❑ **Existing object detection systems include:**
 - DPM
 - This method uses a sliding window technique, where a classifier is applied at regular intervals across the image.
 - R-CNN
 - This approach generates potential bounding boxes using region proposals and then applies a classifier to these boxes.
 - ❑ However, both DPM and R-CNN involve intricate multi-stage processes, making them relatively slow and difficult to fine-tune **as each individual component requires separate training**



Introduction (2/3)

- ❑ **PROBLEM:** It is difficult to emulate the human eye's ability to detect objects instantly and effortlessly.
 - But we need to solve this fundamental challenge in computer vision
- ❑ Typical object detection workflows include:
 - Extracting meaningful features from the input image.
 - Using classifiers or localizers to detect objects in the feature space.
 - Running these models across the full image or selected regions, often using a sliding window approach.
- ❑ The layered nature of these pipelines makes object detection systems slow and unsuitable for real-time applications.



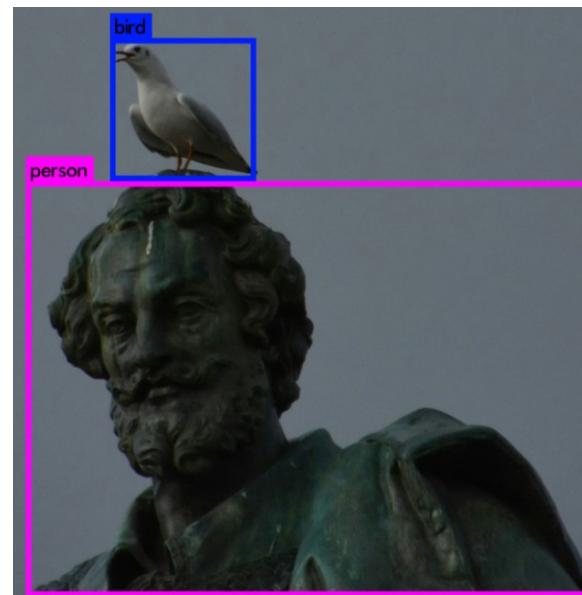
❑ YOLO Model

- Transforms object detection into a unified regression task.
- Image pixels => Bounding boxes => Class probabilities of boxes.
- Utilizes a single NN to predict multiple bounding boxes and their class probabilities in one pass.
- Designed to be fast and simple, avoiding multi-step pipelines.

Introduction (3/3)

❑ YOLO vs DPM vs R-CNN: Advantage

- Extremely fast, capable of handling real-time video streams with only ~ 25 ms delay.
- Learns features that generalize well.
- Makes predictions by considering the entire image context, unlike localized approaches of DPM and R-CNN.



YOLO Unified Detection

□ YOLO splits the input image into an $S \times S$ grid

- Each cell in the grid predicts B bounding boxes along with a confidence score.
- The confidence score reflects two things:
 1. The likelihood that an object exists in the box.
 2. How accurately the predicted box overlaps with the actual object
 - **Confidence = $\text{Pr}(\text{Object}) * \text{IOU}_{\text{Pred}}^{\text{truth}}$**

□ Each bounding box includes 5 predicted values: $x, y, w, h, \text{and } \text{confidence}$

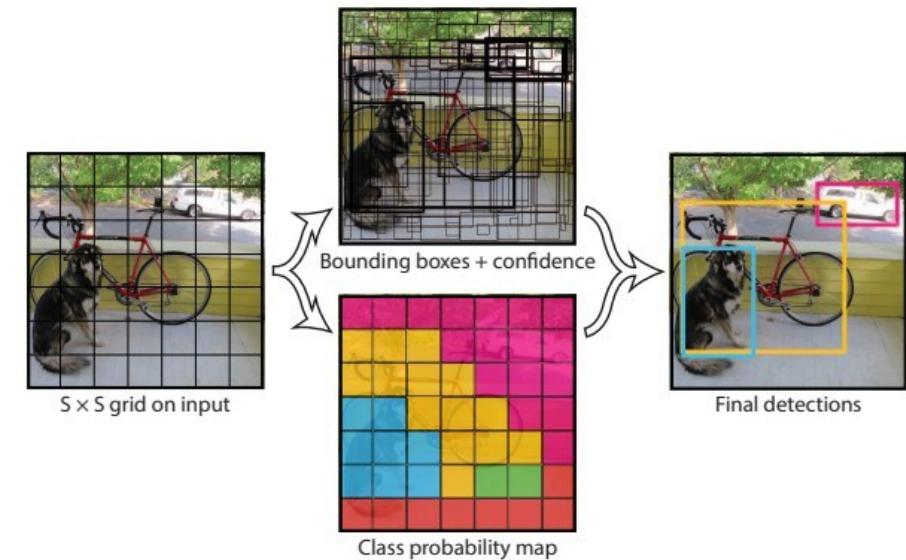
- (x, y) : center of the box relative to the grid cell.
- w and h are relative to the overall image size.
- Confidence prediction is the IOU between the predicted box and ground truth

□ Each grid cell also predicts C conditional class probabilities

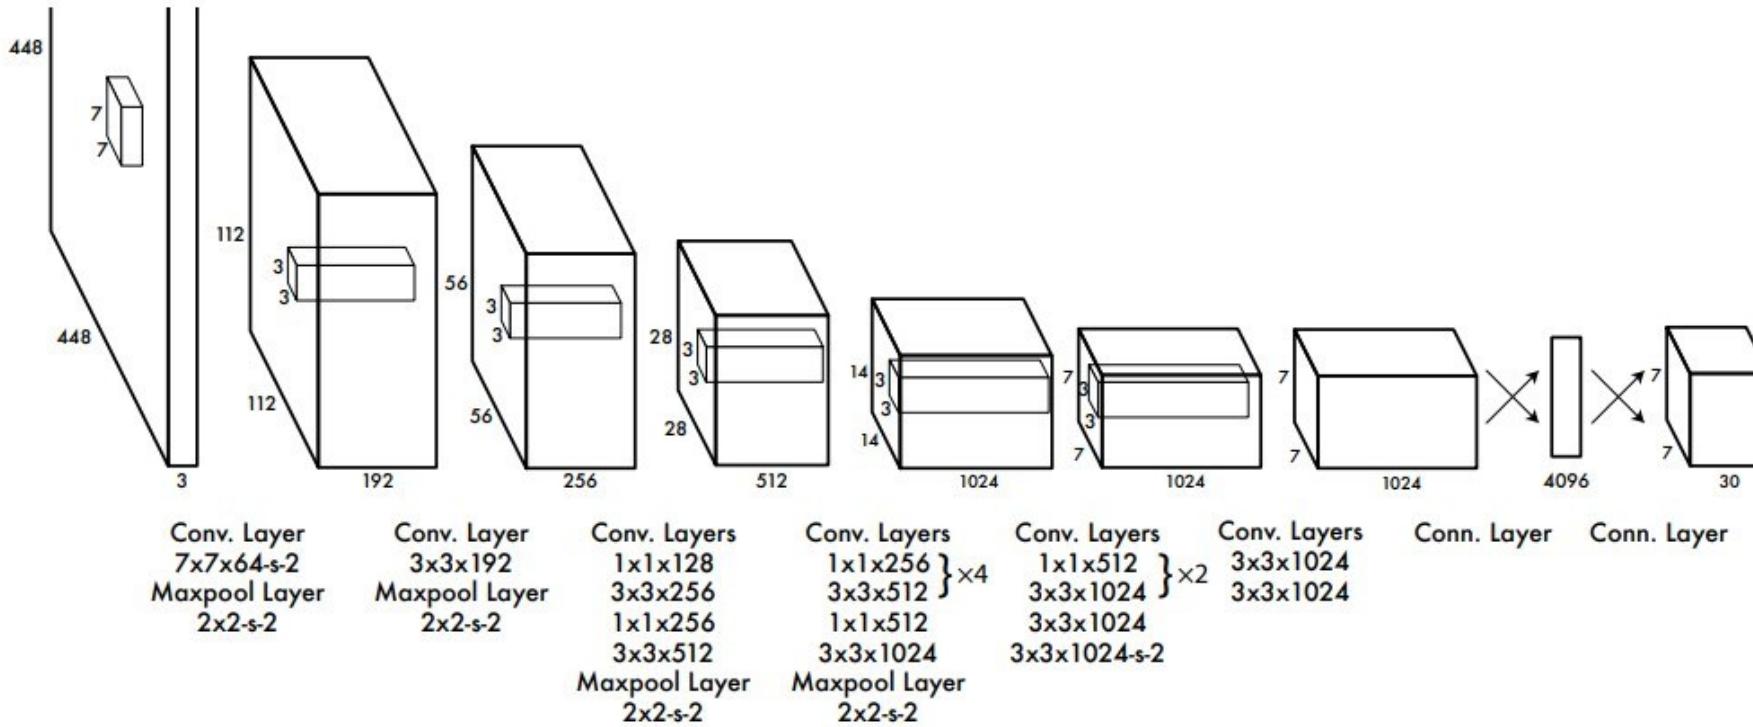
The probabilities are conditioned on the grid cell containing an object

- These are the probabilities of a class given an object is present $\text{Pr}(\text{Class}_i | \text{Object})$.
- Class prediction occurs only if the grid cell contains an object.
- Each class-specific confidence score combines:
 1. Probability of the class appearing in the box.
 2. How well the box aligns with the object it is supposed to detect.

$$\text{Pr}(\text{Class}_i | \text{Object}) * \text{Pr}(\text{object}) * \text{IOU}_{\text{pred}}^{\text{truth}} = \text{Pr}(\text{Class}_i * \text{IOU}_{\text{pred}}^{\text{truth}})$$



YOLO Network Design



- The architecture uses 24 conv layers to extract image features, followed by 2 fully connected layers that predict class probabilities and bounding box coordinates.
- Fast YOLO reduces complexity by using 9 conv layers and a smaller number of filters.

Training & Testing YOLO (1/5)

loss function:

$$\begin{aligned} & \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 \right] \\ & + \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(\sqrt{w_i} - \sqrt{\hat{w}_i})^2 + (\sqrt{h_i} - \sqrt{\hat{h}_i})^2 \right] \\ & + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} (C_i - \hat{C}_i)^2 \\ & + \lambda_{\text{noobj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{noobj}} (C_i - \hat{C}_i)^2 \\ & + \sum_{i=0}^{S^2} \mathbb{1}_i^{\text{obj}} \sum_{c \in \text{classes}} (p_i(c) - \hat{p}_i(c))^2 \quad (3) \end{aligned}$$

where $\mathbb{1}_i^{\text{obj}}$ denotes if object appears in cell i and $\mathbb{1}_{ij}^{\text{obj}}$ denotes that the j th bounding box predictor in cell i is “responsible” for that prediction.

λ_{coord} = grid cell contains an object
 λ_{noobj} = grid cell contains no object



Optimized during training:

- Model was trained for 135 epochs .
- Learning rate rises from 10^{-3} to 10^{-2} .

- The Model is trained on the Pascal Visual Object Classes (VOC) data
 - Darknet framework is used for all training and inference

Training & Testing YOLO (2/5)

```
int convolutional_out_height(convolutional_layer l)
{
    return (l.h + 2*l.pad - l.size) / l.stride + 1;
}

int convolutional_out_width(convolutional_layer l)
{
    return (l.w + 2*l.pad - l.size) / l.stride + 1;

image get_convolutional_image(convolutional_layer l)
{
    return float_to_image(l.out_w,l.out_h,l.out_c,l.output);
}

image get_convolutional_delta(convolutional_layer l)
{
    return float_to_image(l.out_w,l.out_h,l.out_c,l.delta);
}
```



- Darknet C implementation of how to compute the height and width of an output feature map after applying the convolutional layer in terms of **Filter size**, **Stride** and **Padding**
- returns an image representation of the output feature map from the convolutional layer.
- Returns the gradient (delta) of the convolutional layer as an image.

Training & Testing YOLO (3/5)

```
void forward_maxpool_layer(const maxpool_layer l, network net)
{
    int b,i,j,k,m,n;
    int w_offset = -l.pad/2;
    int h_offset = -l.pad/2;

    int h = l.out_h;
    int w = l.out_w;
    int c = l.c;

    for(b = 0; b < l.batch; ++b){
        for(k = 0; k < c; ++k){
            for(i = 0; i < h; ++i){
                for(j = 0; j < w; ++j){
                    int out_index = j + w*(i + h*(k + c*b));
                    float max = -FLT_MAX;
                    int max_i = -1;
                    for(n = 0; n < l.size; ++n){
                        for(m = 0; m < l.size; ++m){
                            int cur_h = h_offset + i*l.stride + n;
                            int cur_w = w_offset + j*l.stride + m;
                            int index = cur_w + l.w*(cur_h + l.h*(k + b*l.c));
                            int valid = (cur_h >= 0 && cur_h < l.h &&
                                         cur_w >= 0 && cur_w < l.w);
                            float val = (valid != 0) ? net.input[index] : -FLT_MAX;
                            max_i = (val > max) ? index : max_i;
                            max = (val > max) ? val : max;
                        }
                    }
                    l.output[out_index] = max;
                    l.indexes[out_index] = max_i;
                }
            }
        }
    }
}

void backward_maxpool_layer(const maxpool_layer l, network net)
{
    int i;
    int h = l.out_h;
    int w = l.out_w;
    int c = l.c;
    for(i = 0; i < h*w*c*l.batch; ++i){
        int index = l.indexes[i];
        net.delta[index] += l.delta[i];
    }
}
```



- Darknet C implementation of forward and backward passes of a Max Pooling Layer

Training & Testing YOLO (4/5)

```

layer make_connected_layer(int batch, int inputs, int outputs, ACTIVATION activation, int batch_normalize, int adam)
{
    int i;
    layer l = {0};
    l.learning_rate_scale = 1;
    l.type = CONNECTED;

    l.inputs = inputs;
    l.outputs = outputs;
    l.batch=batch;
    l.batch_normalize = batch_normalize;
    l.h = 1;
    l.w = 1;
    l.c = inputs;
    l.out_h = 1;
    l.out_w = 1;
    l.out_c = outputs;

    l.output = calloc(batch*outputs, sizeof(float));
    l.delta = calloc(batch*outputs, sizeof(float));

    l.weight_updates = calloc(inputs*outputs, sizeof(float));
    l.bias_updates = calloc(outputs, sizeof(float));

    l.weights = calloc(outputs*inputs, sizeof(float));
    l.biases = calloc(outputs, sizeof(float));

    l.forward = forward_connected_layer;
    l.backward = backward_connected_layer;
    l.update = update_connected_layer;

    //float scale = 1./sqrt(inputs);
    float scale = sqrt(2./inputs);
    for(i = 0; i < outputs*inputs; ++i){
        l.weights[i] = scale*rand_uniform(-1, 1);
    }

    for(i = 0; i < outputs; ++i){
        l.biases[i] = 0;
    }

    if(adam){
        l.m = calloc(l.inputs*l.outputs, sizeof(float));
        l.v = calloc(l.inputs*l.outputs, sizeof(float));
        l.bias_m = calloc(l.outputs, sizeof(float));
        l.scale_m = calloc(l.outputs, sizeof(float));
        l.bias_v = calloc(l.outputs, sizeof(float));
        l.scale_v = calloc(l.outputs, sizeof(float));
    }
    if(batch_normalize){
        l.scales = calloc(outputs, sizeof(float));
        l.scale_updates = calloc(outputs, sizeof(float));
        for(i = 0; i < outputs; ++i){
            l.scales[i] = 1;
        }

        l.mean = calloc(outputs, sizeof(float));
        l.mean_delta = calloc(outputs, sizeof(float));
        l.variance = calloc(outputs, sizeof(float));
        l.variance_delta = calloc(outputs, sizeof(float));

        l.rolling_mean = calloc(outputs, sizeof(float));
        l.rolling_variance = calloc(outputs, sizeof(float));

        l.x = calloc(batch*outputs, sizeof(float));
        l.x_norm = calloc(batch*outputs, sizeof(float));
    }
}

```

```

void forward_connected_layer(layer l, network net)
{
    fill_cpu(l.outputs*l.batch, 0, l.output, 1);
    int m = l.batch;
    int k = l.inputs;
    int n = l.outputs;
    float *a = net.input;
    float *b = l.weights;
    float *c = l.output;
    gemm(0,1,m,n,k,1,a,k,b,k,1,c,n);
    if(l.batch_normalize){
        forward_batchnorm_layer(l, net);
    } else {
        add_bias(l.output, l.biases, l.batch, l.outputs, 1);
    }
    activate_array(l.output, l.outputs*l.batch, l.activation);
}

void backward_connected_layer(layer l, network net)
{
    gradient_array(l.output, l.outputs*l.batch, l.activation, l.delta);

    if(l.batch_normalize){
        backward_batchnorm_layer(l, net);
    } else {
        backward_bias(l.bias_updates, l.delta, l.batch, l.outputs, 1);
    }

    int m = l.outputs;
    int k = l.batch;
    int n = l.inputs;
    float *a = l.delta;
    float *b = net.input;
    float *c = l.weight_updates;
    gemm(1,0,m,n,k,1,a,M,b,0,1,c,n);

    m = l.batch;
    k = l.outputs;
    n = l.inputs;

    a = l.delta;
    b = l.weights;
    c = net.delta;

    if(c) gemm(0,0,m,n,k,1,a,k,b,n,1,c,n);
}

```

```

void forward_connected_layer_gpu(layer l, network net)
{
    fill_gpu(l.outputs*l.batch, 0, l.output_gpu, 1);

    int m = l.batch;
    int k = l.inputs;
    int n = l.outputs;
    float *a = net.input_gpu;
    float *b = l.weights_gpu;
    float *c = l.output_gpu;
    gemm_gpu(0,1,m,n,k,1,a,k,b,k,1,c,n);

    if (l.batch_normalize) {
        forward_batchnorm_layer_gpu(l, net);
    } else {
        add_bias_gpu(l.output_gpu, l.biases_gpu, l.batch, l.outputs, 1);
    }
    activate_array_gpu(l.output_gpu, l.outputs*l.batch, l.activation);
}

void backward_connected_layer_gpu(layer l, network net)
{
    constrain_gpu(l.outputs*l.batch, 1, l.delta_gpu, 1);
    gradient_array_gpu(l.output_gpu, l.outputs*l.batch, l.activation, l.delta_gpu);
    if(l.batch_normalize){
        backward_batchnorm_layer_gpu(l, net);
    } else {
        backward_bias_gpu(l.bias_updates_gpu, l.delta_gpu, l.batch, l.outputs, 1);
    }

    int m = l.outputs;
    int k = l.batch;
    int n = l.inputs;
    float *a = l.delta_gpu;
    float *b = net.input_gpu;
    float *c = l.weight_updates_gpu;
    gemm_gpu(1,0,m,n,k,1,a,M,b,0,1,c,n);

    m = l.batch;
    k = l.outputs;
    n = l.inputs;

    a = l.delta_gpu;
    b = l.weights_gpu;
    c = net.delta_gpu;

    if(c) gemm_gpu(0,0,m,n,k,1,a,k,b,n,1,c,n);
}
#endif

```

Training & Testing YOLO (5/5)



- Testing pretrained YOLO model on images

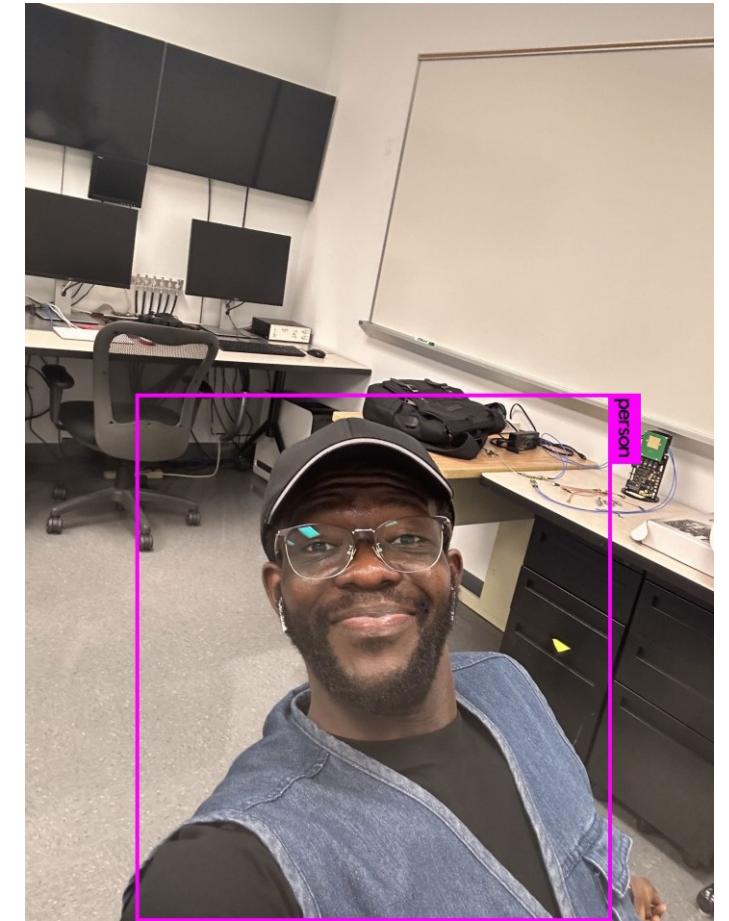
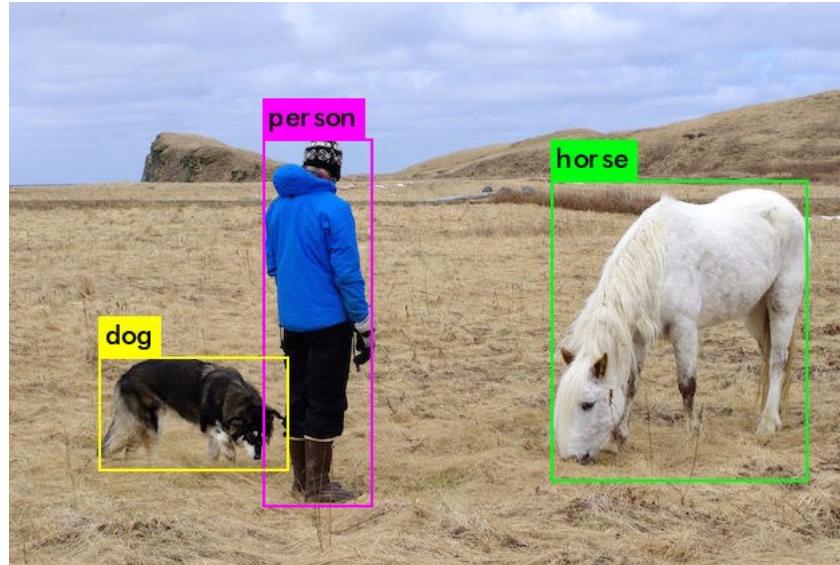
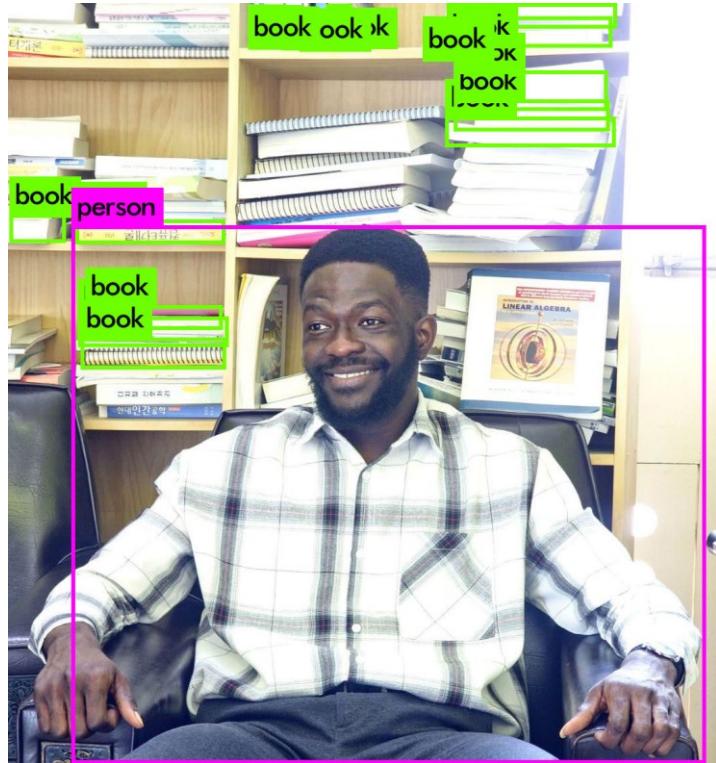
layer	filters	size	input	output
0 conv	32	3 x 3 / 1	608 x 608 x 3 ->	608 x 608 x 32 0.639 BFLOPs
1 conv	64	3 x 3 / 2	608 x 608 x 32 ->	304 x 304 x 64 3.407 BFLOPs
2 conv	32	1 x 1 / 1	304 x 304 x 64 ->	304 x 304 x 32 0.379 BFLOPs
3 conv	64	3 x 3 / 1	304 x 304 x 32 ->	304 x 304 x 64 3.407 BFLOPs
4 res	1		304 x 304 x 64 ->	304 x 304 x 64
5 conv	128	3 x 3 / 2	304 x 304 x 64 ->	152 x 152 x 128 3.407 BFLOPs
6 conv	64	1 x 1 / 1	152 x 152 x 128 ->	152 x 152 x 64 0.379 BFLOPs
7 conv	128	3 x 3 / 1	152 x 152 x 64 ->	152 x 152 x 128 3.407 BFLOPs
8 res	5		152 x 152 x 128 ->	152 x 152 x 128
9 conv	64	1 x 1 / 1	152 x 152 x 128 ->	152 x 152 x 64 0.379 BFLOPs
10 conv	128	3 x 3 / 1	152 x 152 x 64 ->	152 x 152 x 128 3.407 BFLOPs
11 res	8		152 x 152 x 128 ->	152 x 152 x 128
12 conv	256	3 x 3 / 2	152 x 152 x 128 ->	76 x 76 x 256 3.407 BFLOPs
13 conv	128	1 x 1 / 1	76 x 76 x 256 ->	76 x 76 x 128 0.379 BFLOPs
14 conv	256	3 x 3 / 1	76 x 76 x 128 ->	76 x 76 x 256 3.407 BFLOPs
15 res	12		76 x 76 x 256 ->	76 x 76 x 256
16 conv	128	1 x 1 / 1	76 x 76 x 256 ->	76 x 76 x 128 0.379 BFLOPs
17 conv	256	3 x 3 / 1	76 x 76 x 128 ->	76 x 76 x 256 3.407 BFLOPs
18 res	15		76 x 76 x 256 ->	76 x 76 x 256
19 conv	128	1 x 1 / 1	76 x 76 x 256 ->	76 x 76 x 128 0.379 BFLOPs
20 conv	256	3 x 3 / 1	76 x 76 x 128 ->	76 x 76 x 256 3.407 BFLOPs
21 res	18		76 x 76 x 256 ->	76 x 76 x 256
22 conv	128	1 x 1 / 1	76 x 76 x 256 ->	76 x 76 x 128 0.379 BFLOPs
23 conv	256	3 x 3 / 1	76 x 76 x 128 ->	76 x 76 x 256 3.407 BFLOPs
24 res	21		76 x 76 x 256 ->	76 x 76 x 256
25 conv	128	1 x 1 / 1	76 x 76 x 256 ->	76 x 76 x 128 0.379 BFLOPs
26 conv	256	3 x 3 / 1	76 x 76 x 128 ->	76 x 76 x 256 3.407 BFLOPs
27 res	24		76 x 76 x 256 ->	76 x 76 x 256
28 conv	128	1 x 1 / 1	76 x 76 x 256 ->	76 x 76 x 128 0.379 BFLOPs
29 conv	256	3 x 3 / 1	76 x 76 x 128 ->	76 x 76 x 256 3.407 BFLOPs
30 res	27		76 x 76 x 256 ->	76 x 76 x 256
31 conv	128	1 x 1 / 1	76 x 76 x 256 ->	76 x 76 x 128 0.379 BFLOPs
32 conv	256	3 x 3 / 1	76 x 76 x 128 ->	76 x 76 x 256 3.407 BFLOPs
33 res	30		76 x 76 x 256 ->	76 x 76 x 256
34 conv	128	1 x 1 / 1	76 x 76 x 256 ->	76 x 76 x 128 0.379 BFLOPs
35 conv	256	3 x 3 / 1	76 x 76 x 128 ->	76 x 76 x 256 3.407 BFLOPs
36 res	33		76 x 76 x 256 ->	76 x 76 x 256
37 conv	512	3 x 3 / 2	76 x 76 x 256 ->	38 x 38 x 512 3.407 BFLOPs
38 conv	256	1 x 1 / 1	38 x 38 x 512 ->	38 x 38 x 256 0.379 BFLOPs
39 conv	512	3 x 3 / 1	38 x 38 x 256 ->	38 x 38 x 512 3.407 BFLOPs
40 res	37		38 x 38 x 512 ->	38 x 38 x 512
41 conv	256	1 x 1 / 1	38 x 38 x 512 ->	38 x 38 x 256 0.379 BFLOPs
42 conv	512	3 x 3 / 1	38 x 38 x 256 ->	38 x 38 x 512 3.407 BFLOPs
43 res	40		38 x 38 x 512 ->	38 x 38 x 512
44 conv	256	1 x 1 / 1	38 x 38 x 512 ->	38 x 38 x 256 0.379 BFLOPs
45 conv	512	3 x 3 / 1	38 x 38 x 256 ->	38 x 38 x 512 3.407 BFLOPs
46 res	43		38 x 38 x 512 ->	38 x 38 x 512
47 conv	256	1 x 1 / 1	38 x 38 x 512 ->	38 x 38 x 256 0.379 BFLOPs
48 conv	512	3 x 3 / 1	38 x 38 x 256 ->	38 x 38 x 512 3.407 BFLOPs
49 res	46		38 x 38 x 512 ->	38 x 38 x 512
50 conv	256	1 x 1 / 1	38 x 38 x 512 ->	38 x 38 x 256 0.379 BFLOPs
51 conv	512	3 x 3 / 1	38 x 38 x 256 ->	38 x 38 x 512 3.407 BFLOPs
52 res	49		38 x 38 x 512 ->	38 x 38 x 512
53 conv	256	1 x 1 / 1	38 x 38 x 512 ->	38 x 38 x 256 0.379 BFLOPs
54 conv	512	3 x 3 / 1	38 x 38 x 256 ->	38 x 38 x 512 3.407 BFLOPs
55 res	52		38 x 38 x 512 ->	38 x 38 x 512
56 conv	256	1 x 1 / 1	38 x 38 x 512 ->	38 x 38 x 256 0.379 BFLOPs
57 conv	512	3 x 3 / 1	38 x 38 x 256 ->	38 x 38 x 512 3.407 BFLOPs
58 res	55		38 x 38 x 512 ->	38 x 38 x 512
59 conv	256	1 x 1 / 1	38 x 38 x 512 ->	38 x 38 x 256 0.379 BFLOPs
60 conv	512	3 x 3 / 1	38 x 38 x 256 ->	38 x 38 x 512 3.407 BFLOPs

Loading weights from yolov3.weights...Done!
data/horses.jpg: Predicted in 6.276225 seconds.
horse: 100%
horse: 100%
horse: 96%
horse: 96%
horse: 95%

55 res 52
56 conv 256 1 x 1 / 1
57 conv 512 3 x 3 / 1
58 res 55
59 conv 256 1 x 1 / 1
60 conv 512 3 x 3 / 1
61 res 58
62 conv 1024 3 x 3 / 2
63 conv 512 1 x 1 / 1
64 conv 1024 3 x 3 / 1
65 res 62
66 conv 512 1 x 1 / 1
67 conv 1024 3 x 3 / 1
68 res 65
69 conv 512 1 x 1 / 1
70 conv 1024 3 x 3 / 1
71 res 68
72 conv 512 1 x 1 / 1
73 conv 1024 3 x 3 / 1
74 res 71
75 conv 512 1 x 1 / 1
76 conv 1024 3 x 3 / 1
77 conv 512 1 x 1 / 1
78 conv 1024 3 x 3 / 1
79 conv 512 1 x 1 / 1
80 conv 1024 3 x 3 / 1
81 conv 255 1 x 1 / 1
82 yolo
83 route 79
84 conv 256 1 x 1 / 1
85 upsample 2x
86 route 85 61
87 conv 256 1 x 1 / 1
88 conv 256 1 x 1 / 1
89 conv 256 1 x 1 / 1
90 conv 512 3 x 3 / 1
91 conv 256 1 x 1 / 1
92 conv 512 3 x 3 / 1
93 conv 255 1 x 1 / 1
94 yolo
95 route 91
96 conv 128 1 x 1 / 1
97 upsample 2x
98 route 97 36
99 conv 128 1 x 1 / 1
100 conv 256 3 x 3 / 1
101 conv 128 1 x 1 / 1
102 conv 256 3 x 3 / 1
103 conv 128 1 x 1 / 1
104 conv 256 3 x 3 / 1
105 conv 255 1 x 1 / 1
106 yolo

Loading weights from yolov3.weights...Done!
data/eddy_1.jpg: Predicted in 6.005917 seconds.
book: 76%
book: 65%
book: 64%
book: 64%
book: 62%
book: 62%
book: 60%
book: 60%
book: 59%
book: 58%
book: 56%
book: 52%
book: 50%
person: 99%

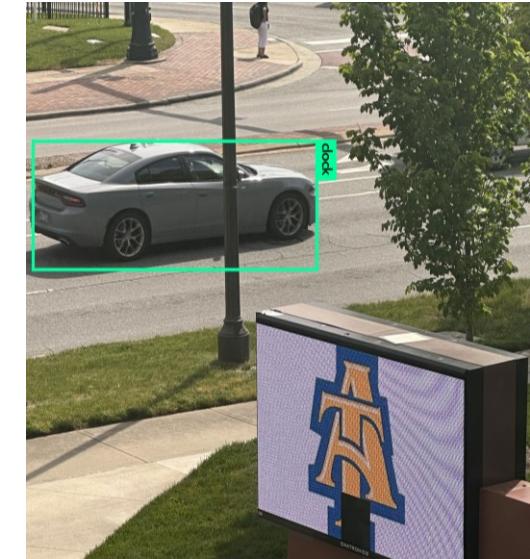
YOLO Result



- Testing pretrained YOLO model on the image

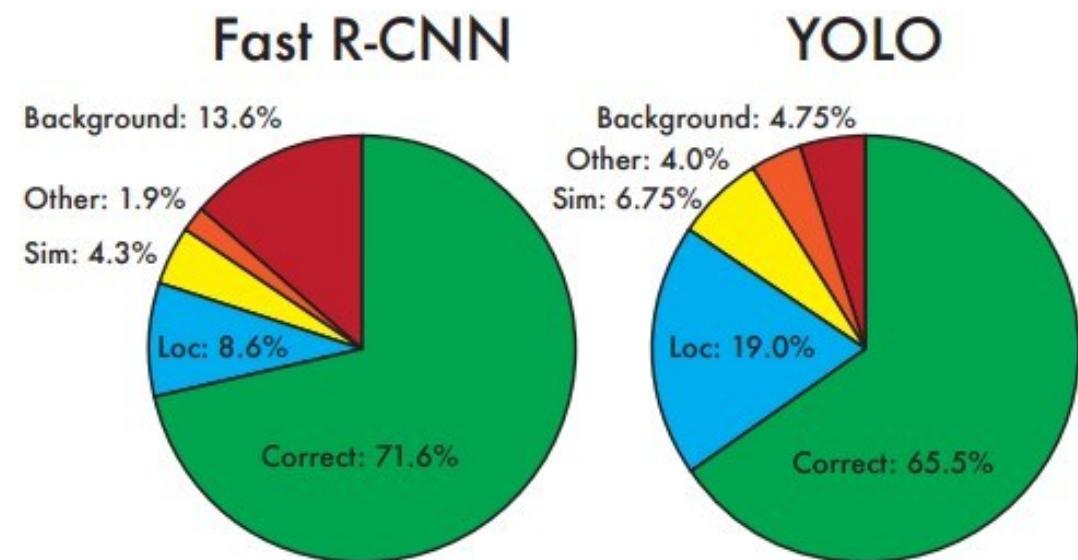
YOLO Limitations

- ❑ Because YOLO learns to estimate bounding boxes directly from the data, it has difficulty adapting to objects with unfamiliar or unusual shapes.
- ❑ The loss function does not distinguish between errors in small and large bounding boxes:
 - A minor error in a large bounding box is often negligible, but a similar error in a small box significantly impacts the IOU score.
 - Most of the errors come from inaccurate object localization.
- ❑ Struggles with precision – While it quickly detects objects, it sometimes fails to accurately pinpoint their locations.
- ❑ YOLO places strict spatial constraints, which reduces its ability to detect multiple closely located objects.

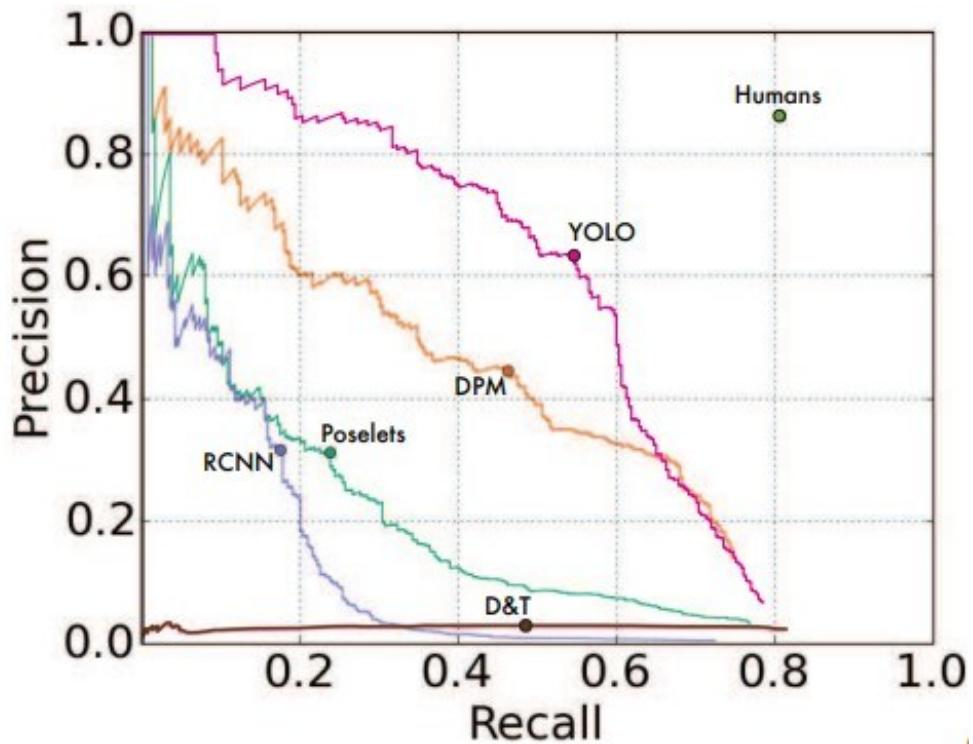


YOLO vs Other Object Detection Systems (1/2)

Real-Time Detectors	Train	mAP	FPS
100Hz DPM [30]	2007	16.0	100
30Hz DPM [30]	2007	26.1	30
Fast YOLO	2007+2012	52.7	155
YOLO	2007+2012	63.4	45
Less Than Real-Time			
Fastest DPM [37]	2007	30.4	15
R-CNN Minus R [20]	2007	53.5	6
Fast R-CNN [14]	2007+2012	70.0	0.5
Faster R-CNN VGG-16[27]	2007+2012	73.2	7
Faster R-CNN ZF [27]	2007+2012	62.1	18
YOLO VGG-16	2007+2012	66.4	21



YOLO vs Other Object Detection Systems (2/2)



	VOC 2007 AP	Picasso		People-Art AP
	AP	Best F_1		
YOLO	59.2	53.3	0.590	45
R-CNN	54.2	10.4	0.226	26
DPM	43.2	37.8	0.458	32
Poselets [2]	36.5	17.8	0.271	
D&T [4]	-	1.9	0.051	



Conclusion

- ❑ YOLO introduces a unified framework for object detection:
 - It's easy to build and can be trained on entire images.
 - Unlike traditional methods that treat detection as a classification problem, YOLO uses a loss function tailored for detection accuracy and trains the model end-to-end.
- ❑ YOLO adapts effectively to new environments, making it well-suited for real-time object detection tasks that require speed and reliability.
- ❑ Fast YOLO is recognized as the quickest general-purpose object detection model available and is capable of real-time detection.



Any Questions?

E-mail: [ekwao at aggies.ncat.edu](mailto:ekwao@aggies.ncat.edu)